

AI Lab - Session 6

Deep Reinforcement Learning

Davide Corsi
Alessandro Farinelli

University of Verona
Department of Computer Science

January 17th 2022



UNIVERSITÀ
di **VERONA**

Dipartimento
di **INFORMATICA**

Start Your Working Environment

Start the previously installed (Session 1) conda environment *ai-lab*

Listing 1: Update Environment

```
cd AI-Lab
git stash (NB: remember to backup the previous lessons before this step!)
git pull
git stash pop
conda activate ai-lab
pip install tensorflow
pip install keras
jupyter notebook
```

Listing 2: Open Lesson

To open the tutorial navigate with your browser to:
`lesson_5/lesson_5_problem.ipynb`

What is it

Keras is a high-level neural networks APIs. It is written in Python and supports multiple back-end neural network computation engines:

- *built on top of TensorFlow 2.0*
- *optimized to work both on CPU and GPU*
- *simple functions to create, train and modify neural networks with state of the art architecture*

What is it for

Designed to enable fast experimentation with deep neural networks, it focuses on being user-friendly, modular, and extensible library.

Where to find it

<https://keras.io/>

Assignments

- Your assignments for this session are:
lesson_5/lesson_5_problem.ipynb. You will be required to implement Deep Reinforcement Learning Algorithms, in particular the main loop and the "train" function
- In the following you can find the pseudocode

Deep Q-Learning

Input: *environment, neural_network, trials, expl_param, score_queue*

Output: *neural_network, score_queue*

```
1: initialize the experience buffer
2: initialize the score queue
3: for  $i \leftarrow 0$  to trials do
4:     initialize s observe current state
5:     repeat
6:         Select and execute action a
7:         Observe new state  $s'$  and receive immediate reward r
8:         Add (s, a,  $s'$ , r) to experience buffer
9:         TRAIN_FUNC(neural_network, experience_buffer)
10:        update state  $s \leftarrow s'$ 
11:    until s is terminal
12:    update score_queue
13:    if score_queue[i] > goal_score then
14:        break loop
15: return neural_network, score_queue
```

▷ A fixed size queue
▷ An infinite size queue
▷ ϵ -greedy approach

Train Function

Input: *neural_network*, *experience_buffer*(*MB*), *gamma*

Output: *neural_network*

```
1: Sample mini-batch MB of experiences from buffer
2: for  $s, a, s', r \in \text{MB}$  do                                ▷ (state, action, next_state, reward)
3:    $\text{target} \leftarrow \text{PREDICT}(\text{neural\_network}, s)$ 
4:   if  $s'$  is terminal then
5:      $\text{target}[a] = r$ 
6:   else
7:      $\text{max-q} = \max(\text{PREDICT}(\text{neural\_network}, s'))$           ▷ max q-value from  $s'$ 
8:      $\text{target}[a] = r + (\text{max-q} * \text{gamma})$ 
9:    $\text{neural\_network} \leftarrow \text{FIT}(\text{neural\_network}, s, \text{target})$     ▷ back-propagation
10: return neural_network
```